

```

## STAT LEARNING 2
## HOMEWORK 4, Problem 5

# Read in necessary Libraries
library(ggplot2)
library(magrittr)
library(R.matlab)
library(geometry)
library(dplyr)

# Set the working directory
setwd("C:/Users/jrdha/OneDrive/Desktop/USU_Fa2018/Moon__SLDM2/hw4/Problem5")

# This function kicks off the recursive subGradient function.
# It takes as arguments: max_Iter (max number of iterations), predictors
# (dataframe of input features), response (dataframe containing values of
# responses), lambda (used for calculation of regularization term).
# The function returns optimized theta vector (weights, b).
initSubGradient <- function(max_Iter, predictors, response, lambda){
  b = 1
  ws = rep(0, ncol(predictors) - 1)
  inittheta <- c(b, ws)

  # Dataframe to store info from iterations
  iteration_Info <- data.frame(numItr = double(0), objFun = double(0),
                               w1 = double(0), w2 = double(0), b = double(0))

  # call subGradient function
  theta <- subGradient(theta = inittheta, num_Iter = 1, max_Iter = max_Iter,
                       xs = predictors, y = response,
                       lambda = lambda, iteration_Info = iteration_Info)
  return(theta)
}

# This function (recursively) performs the calculations to determine optimal value
# of theta or calls itself again.
# the initSubGradient function.
# Takes as arguments: theta (vector of weights and b), num_Iter and max_Iter (are
# self-explanatory), xs (dataframe containing input data), y (dataframe containing
# response data), lambda (regularization parameter), iteration_Info (dataframe
# that stores parameter estimates and value of objective function).
subGradient <- function(theta, num_Iter, max_Iter, xs, y, lambda, iteration_Info){

  # Alpha is the step size, n is the number of instances in preds
  alpha <- 100/num_Iter
  n <- nrow(preds)

```

```

# Either return optimized theta and iteration info, or keep going
if(num_Iter > max_Iter){
  return(list("theta" = theta, "iteration_Info" = iteration_Info))
} else {
  print(paste("Iteration: ", num_Iter))

  u <- double(length(theta))
  for(i in 1:nrow(xs)){
    # calculate gradient
    grad <- calcGrad(b = theta[1], w = theta[2:length(theta)],
                     yi = as.numeric(y$Y[i]), xi = as.numeric(xs[i,]),
                     n = n, lambda = lambda)

    # Update u
    u <- u + grad
  }

  # Having calculated gradient, take a step in the opposite direction
  theta.1 <- theta - (alpha * u)

  # calculate value of objective function
  objFctn_val <- calc_objFctn(preds = xs, resp = y, theta = theta.1, lambda = lambda)

  # This row will be added to the bottom of the iteration_Info dataframe
  hnew <- data.frame(numItr = num_Iter, objFun = objFctn_val, b = theta.1[1],
                    w1 = theta.1[2], w2 = theta.1[3])

  iteration_Info <- rbind(iteration_Info, hnew)
  num_Iter <- num_Iter + 1

  return(subGradient(theta = theta.1, num_Iter = num_Iter, max_Iter = max_Iter,
                    xs = xs, y = y, lambda = lambda, iteration_Info = iteration_Info
  ))
}
}

```

```

# This function kicks off the recursive stochSubGradient function.
# It takes as arguments: max_Iter, predictors, response, and lambda (we've seen
# all of these before, so I won't re-explain what they are).
# It returns the optimal value of the theta vector.
initStochSubGradient <- function(max_Iter, predictors, response, lambda){

```

```

  # Initial guesses for theta (b and the weights)
  b = 1
  ws = rep(0, ncol(predictors) - 1)
  inittheta <- c(b, ws)

```

```

  # Initialize the iteration_Info data frame
  iteration_Info <- data.frame(numItr = double(0), objFun = double(0),
                              w1 = double(0), w2 = double(0), b = double(0))

```

```

# Recursion of stochSubGradient() begins
theta <- stochSubGradient(theta = inittheta, num_Iter = 1, max_Iter = max_Iter,
                          xs = predictors, y = response,
                          lambda = lambda, iteration_Info = iteration_Info)

return(theta)
}

# This recursive function calls the subGradient function.
# Takes as arguments all the things we know and love: theta vector, num_Iter,
# max_Iter, xs, y, lambda, and iteration_Info dataframe.
# This function either returns the optimal value of theta, or calls itself again
# depending on how many iterations have taken place.
stochSubGradient <- function(theta, num_Iter, max_Iter, xs, y, lambda, iteration_In
fo){

  # Set the step size, how many predictor instances there are, and recommended
  # minibatch size
  alpha <- 100/num_Iter
  n <- nrow(preds)
  m <- 1

  # The use of the sample() function below will be how we randomly sample
  if(num_Iter > max_Iter){
    return(list("theta" = theta, "iteration_Info" = iteration_Info))
  } else {
    print(paste("Iteration Number: ", num_Iter))
    index <- 1:n
    rand.index <- sample(index, n)
    for(i in rand.index){
      grad <- calcGrad(b = theta[1], w = theta[2:length(theta)],
                      yi = as.numeric(y$Y[i]), xi = as.numeric(xs[i,]),
                      n = 1, lambda = lambda)
      theta <- theta - (alpha * grad) # Move in opposite direction of subGrad
                                     # to update theta.
    }
    theta.1 <- theta

    # Calculate the (current) value of the objective function.
    objFctn_val <- calc_objFctn(preds = xs, resp = y, theta = theta.1, lambda = lam
bda)

    # This row will be added to the bottom of the iteration_Info dataframe
    hnew <- data.frame(numItr = num_Iter, objFun = objFctn_val, b = theta.1[1],
                      w1 = theta.1[2], w2 = theta.1[3])

    iteration_Info <- rbind(iteration_Info, hnew)
    num_Iter <- num_Iter + 1
    return(subGradient(theta = theta.1, num_Iter = num_Iter, max_Iter = max_Iter, x
s = xs, y = y, lambda = lambda, iteration_Info = iteration_Info))
  }
}

```

```

# This function calculates the gradient.
# Takes as arguments: w (weights vector), b (part of theta, intercept term),
# yi (response), xi (vector containing predictor instance), n (num of instances),
# lambda (same parameter for calculating regularization penalty term).
# Returns the subGradient.

```

```

calcGrad <- function(w, b, yi, xi, n, lambda){

  wvec <- c(0, w)
  term <- 1 - (yi * (dot(wvec, xi) + b))
  if(term >= 0){
    gradJi <- (1/n)*(-yi*xi) + (lambda/n)*wvec
  } else {
    gradJi <- (lambda/n)*wvec
  }
  sg <- (gradJi)
  return(sg)
}

```

```

# Calculates the value of the objective function for given inputs.
# Takes as arguments: preds (predictor data), resp (response value), theta (same
# vector of weights and b), lambda (same regularization penalty parameter).

```

```

calc_objFctn <- function(preds, resp, theta, lambda){

  n <- nrow(preds)
  b <- theta[1]
  w <- theta[2:length(theta)]
  wvec <- c(0, w)
  val_Summation <- 0
  for(i in 1:n){
    yi = as.numeric(resp$Y[i])
    xi = as.numeric(preds[i,])
    m1 <- 0
    m2 <- 1 - yi*(dot(wvec, xi) + b)
    term <- max(m1, m2)
    val_Summation <- val_Summation + term
  }
  val_objFctn <- 1/n * (val_Summation) + (lambda/2)*dot(w,w)
  return(val_objFctn)
}

```

```

# This function plots the values of the objective functions against the iteration
# number for that given value of J.
# Takes as arguments the iteration_Info dataframe for plotting, and a title.
# Output is the plot described above.

```

```

plot_objFctn_vals <- function(iteration_Info, title){
  plot(iteration_Info$numItr, iteration_Info$objFun, xlab = "Iteration Number",
        ylab = "value of Objective Function J", type = 'o',
        main = title)
}

```

```

#####
#== Implementing the defined functions to solve the given problem =====
#####

givenData <- R.matlab::readMat("nuclear.mat") %>% lapply(t) %>% lapply(as_tibble)
colnames(givenData[[1]]) <- sprintf("X_%s", seq(1:ncol(givenData[[1]])))
colnames(givenData[[2]]) <- c("Y")
givenData <- bind_cols(givenData) %>% select(Y, everything())

# givenData <- slice(givenData, 1:300)
preds <- select(givenData, X_1, X_2)
X0 <- rep(1, nrow(givenData))
preds <- cbind(X0, preds)
resp <- select(givenData, Y)
resp$Y <- as.numeric(resp$Y)

#== Results of implementing subGradient method =====

# Seems to flatten out after about 35 iterations, chose a fairly small value of Lam
bda.
subGrad_results <- initSubGradient(max_Iter = 35, predictors = preds, response = re
sp, lambda = 0.001)

plot_objFctn_vals(subGrad_results$iteration_Info, title = "SubGradient Descent Meth
od Results")

b <- subGrad_results$theta[1]
w1 <- subGrad_results$theta[2]
w2 <- subGrad_results$theta[3]

slopeVal <- w1/-w2
interceptVal <- b/-w2

# Gets mad if not converted to the factor data type
givenData$Y <- as.factor(givenData$Y)

ggplot(data = givenData, aes(x = X_1, y = X_2, color = Y)) +
  geom_point() +
  scale_color_manual(values=c("-1" = "turquoise", "1" = "magenta")) +
  geom_abline(slope = slopeVal, intercept = interceptVal, lwd = 1) +
  xlab("Values of X1") +
  ylab("Values of X2") +
  ggtitle("Separation using SubGradient Method")

```

```

#=== Results of implementing stochSubGradient method =====
# This one took a lot more iterations to converge as compared to subGradient method
.
stoch_subGrad_results <- initStochSubGradient(max_Iter = 100, predictors = preds,
                                             response = resp, lambda = 0.001)

plot_objFctn_vals(stoch_subGrad_results$iteration_Info, title = "Stochastic SubGradient Descent Method Results")

b <- stoch_subGrad_results$theta[1]
w1 <- stoch_subGrad_results$theta[2]
w2 <- stoch_subGrad_results$theta[3]

slopeVal <- w1/-w2
interceptVal <- b/-w2

plot(preds$X_1, preds$X_2)
abline(a = interceptVal, b = slopeVal)

# Again, have to change this data type to factor or it'll get mad
givenData$Y <- as.factor(givenData$Y)

ggplot(data = givenData, aes(x = X_1, y = X_2, color = Y)) +
  geom_point() +
  scale_color_manual(values=c("-1" = "turquoise", "1" = "magenta")) +
  geom_abline(slope = slopeVal, intercept = interceptVal, lwd = 1) +
  xlab("Values of X1") +
  ylab("Values of X2") +
  ggtitle("Separation using Stochastic SubGradient Method")

```