

```

## STAT LEARNING 2
## HOMEWORK 4, Problem 6

# Libraries needed
library(caret)
library(magrittr)
library("e1071")

# Set working directory
setwd("C:/Users/jrdha/OneDrive/Desktop/USU_Fa2018/Moon__SLDM2/hw4/Problem6")

#=====
#==== INITIAL STEPS =====
#=====

# Read in the data
parkData <- read.csv("parkinsonsData.csv", header = TRUE)

# Subset, get rid of first column
parkData <- parkData[, -1]

# Split into training and test data sets. Going with the typical 80/20 split.
lengthTrain <- round(nrow(parkData) * 0.8)
lengthTest <- nrow(parkData) - lengthTrain

# Set seed
set.seed(1234)

# Make indices for training data subsetting
train_ind <- sample(seq_len(nrow(parkData)), size = lengthTrain)

# Split up the data
trainData <- parkData[train_ind, ]
testData <- parkData[-train_ind, ]

# Make the response a factor (rather than integer) so the algorithms perform
# classification instead of regression.
trainData$status <- as.factor(trainData$status)

```

```

#####
#=== LOGISTIC REGRESSION ===
#####

# This specifies that we'll be doing 10-fold crossvalidation
ctrl <- trainControl(method = "repeatedcv",
                     number = 10,
                     savePredictions = TRUE)

# Train the model
mod_fit <- train(status ~.,
                 data = trainData,
                 method= "glm",
                 family= binomial(),
                 trControl = ctrl,
                 tuneLength = 10)

# This gives the 10-fold crossvalidated training error
trainError_logReg <- 1 - mod_fit$results$Accuracy
trainError_logReg

## [1] 0.1533929

# Generates the predictions for our testData
testData$pred = predict(mod_fit, newdata=testData)

# Generates a column that says whether or not a test observation was correctly
# classified (1==correct, 0==incorrect)
testData$correct <- with(testData, ifelse(status == pred, 1, 0))

# Generate the test error
totalCorrect <- sum(testData$correct)
testError_logReg <- 1 - (totalCorrect / lengthTest)
testError_logReg

## [1] 0.1794872

# RETURNS A TEST ERROR RATE OF 0.1794872, WHICH IS PRECISELY 7
# MISCLASSIFICATIONS. WE WILL BEAT THIS WITH BOTH KERNEL-VERSIONS OF SVM.

# Remove the "correct" and "pred" columns so that we'll have the original data
# to work with for the SVM portion
testData <- subset(testData, select = -c(correct, pred))

```

```

#=====
#=== SVM =====
#=====

#===== TUNING LINEAR-KERNEL MODEL =====
#=====

# Tuning LINEAR KERNEL: Round 1 (used this to get a set of values that did well)
# Ran it by changing cost = 1.digit*10^(-3:3) where I set digit = 0,1,2,...,9
set.seed(2345)
svm_tune_linear <- tune.svm(status~.,
                           data = trainData,
                           kernel = "linear",
                           cost = 1.9*10^(-3:3))

print(svm_tune_linear)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   1.9
##
## - best performance: 0.14125

#===== RESULTS =====
# Parameter tuning of 'svm':
# - sampling method: 10-fold cross validation
# - best parameters:
#   cost
#   1.9
# - best performance: 0.14125

# Having found 1.9 to be the best value from above (by trying 1.0, 1.1, 1.2,...
# ...,1.9) and seeing that most of these values were around 1.5 - 2.0, I decided
# to just look at values between 1 and 2.
# Tuning LINEAR KERNEL: Round 2 (Looking only at values between 1 and 2.1)
set.seed(2345)
svm_tune_linear <- tune.svm(status~.,
                           data = trainData,
                           kernel = "linear",
                           cost = seq(1,2.1,0.01))

print(svm_tune_linear)

```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   1.52
##
## - best performance: 0.14125

#===== RESULTS =====
# IRONICALLY, THIS COMES UP WITH THE SAME CROSSVALIDATED ACCURACY AS FOR C=1.9,
# SO WE SHOULDN'T EXPECT BETTER RESULTS FOR OUR TEST ERROR BY USING C=1.52 OVER
# C=1.9.
# Parameter tuning of 'svm':
#   - sampling method: 10-fold cross validation
#   - best parameters:
#     cost
#     1.52
#   - best performance: 0.14125

# THIS FITS A FINAL SVM MODEL, USING THE CROSSVALIDATED TUNING PARAMETER 1.52
svm_model <- svm(status ~., trainData, kernel = "linear", cost = 1.52)
testData$pred <- predict(svm_model, testData)
testData$correct <- with(testData, ifelse(testData$status == pred, 1, 0))
# Generate the test error
# Total number of correct classifications
totalCorrect <- sum(testData$correct)
testError_SVM <- 1 - (totalCorrect / lengthTest)
testError_SVM

## [1] 0.05128205

# Remove the "correct" column for the SVM portion
testData <- subset(testData, select = -c(correct, pred))

#===== RESULTS =====
# We get the same error rate (0.05128205) as we did by using C=1.9. However,
# this makes sense, as we're already correctly classifying 37/39 values
# correctly with C=1.9, so there's very little room for improvement. Using
# C=1.52 gives just as good of results.
```

```

# ADDITIONAL TUNING: WIDENING THE WINDOW, LOOKING BETWEEN 0.1-19
# ONCE AGAIN, WE GET A COST IN THE SWEET SPOT OF 1.5 - 1.9, RETURNING 1.6, AND
# THE IDENTICAL 0.14125 CROSS-VALIDATED ERROR RATE
set.seed(2345)
svm_tune_linear <- tune.svm(status~.,
                           data = trainData,
                           kernel = "linear",
                           cost = seq(0.1,19,0.1))

print(svm_tune_linear)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   1.6
##
## - best performance: 0.14125

# Parameter tuning of 'svm':
# - sampling method: 10-fold cross validation
# - best parameters:
#   cost
#   1.6
# - best performance: 0.14125

# THIS FITS A FINAL SVM MODEL, USING THE CROSSVALIDATED TUNING PARAMETER 1.6
svm_model <- svm(status ~., trainData, kernel = "linear", cost = 1.6)
testData$pred <- predict(svm_model, testData)
testData$correct <- with(testData, ifelse(testData$status == pred, 1, 0))
# Generate the test error
# Total number of correct classifications
totalCorrect <- sum(testData$correct)
testError_SVM <- 1 - (totalCorrect / lengthTest)
testError_SVM

## [1] 0.05128205

# Remove the "correct" column for the SVM portion
testData <- subset(testData, select = -c(correct, pred))

# COMMON TEST ERROR RATES:
# 0.05128205 for C values in the range of about 1.5 - 1.9
# 0.07692308 for other C values
# 0.1025641 for other C values
# THIS MAKES SENSE: our test data set has only 39 observations. Therefore, we
# can only achieve a certain level of granularity with our test error. The
# best-tuned values for C miss only 2/39 classifications.

```

```

#===== TUNING GAUSSIAN-KERNEL MODEL =====
#=====

# INITIAL PASS (Leave base values at 10)
set.seed(2345)
svm_tune_gaussian <- tune.svm(status~., data = trainData,
                             cost = 10^(-3:7),
                             gamma = 10^(-5:2))

print(svm_tune_gaussian)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   gamma cost
##   0.1    10
##
## - best performance: 0.03208333

# Gives cost = 10 and gamma = 0.1 as best values, CV error of 0.03208333


# LEFT COST=1.6, SINCE THIS WAS THE BEST VALUE FROM LINEAR-KERNEL SVM
# THIS GIVES POORER RESULTS THAN WHEN WE LET THE COST BE BIGGER (and vary from
# c=1.6, see below)
# CHANGE THE BASE VALUES FOR GAMMA, TRYING 1.0, 1.1, 1.2,...,1.9
set.seed(2345)
svm_tune_gaussian <- tune.svm(status~., data = trainData,
                             cost = 1.6,
                             gamma = 1.9*10^(-3:2))

print(svm_tune_gaussian)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   gamma cost
##   0.19    1.6
##
## - best performance: 0.06375

# Gives cost = 1.6 and gamma = 0.19 as best values with CV error: 0.06375

```

```

# MANUAL GRID SEARCH, TRYING ALL 100 POSSIBLE COMBINATIONS OF (1.digit,1.digit)
# for GAMMA and COST
# CHANGE THE BASE VALUES FOR GAMMA, TRYING 1.0, 1.1, 1.2,...,1.9
# CHANGE THE BASE VALUES FOR COST, TRYING 1.0, 1.1, 1.2,...,1.9
set.seed(2345)
svm_tune_gaussian <- tune.svm(status~., data = trainData,
                             cost = 1.0*10^(-2:2),
                             gamma = seq(0.1,0.19,0.01))
print(svm_tune_gaussian)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   gamma cost
##   0.11    10
##
## - best performance: 0.02583333

# for cost=1.0^power, was consistently picking cost=10, and gamma=0.1digit,
# so I set gamma to only be values of 0.1, 0.11, 0.12,...,0.19, and then tried
# different values for cost (changing base to be 1.0, 1.1, 1.2,...,1.9)
# For these restricted gamma, it was consistently using gamma=0.16, and then cost
# of 12, 13, 14, and giving CV error of 0.03833333.
# However, the best combination turned out to be cost=10, gamma=0.11, giving a
# CV error of 0.02583333!

# TRYING TO NARROW DOWN EVEN FURTHER: IF WE KNOW THAT COST=10 AND GAMMA=0.11
# GIVES THE BEST CV ERROR, LET'S TRY OTHER VALUES CLOSE TO THESE ONES.
# I made the steps of the sequences more granular so we could try more values
# and potentially find better ones.
set.seed(2345)
svm_tune_gaussian <- tune.svm(status~., data = trainData,
                             cost = seq(5,15,0.5),
                             gamma = seq(0.1,0.19,0.005))
print(svm_tune_gaussian)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   gamma cost
##   0.105    8
##
## - best performance: 0.02583333

```

```

# HERE'S WHAT IS RETURNED. As we can see, we're still getting the same CV error
# RATE (0.02583333) WE WERE FOR GAMMA=0.11 AND COST=10.
# SO I THINK WE CAN SAFELY CONCLUDE THAT A COST IN 8-10 IS GOOD, AND TUNED
# GAMMA OF EITHER 0.105 OR 0.11 IS GOING TO BE OUR BEST BEST FOR THE
# GAUSSIAN-KERNEL SVM
# Parameter tuning of 'svm':
#   - sampling method: 10-fold cross validation
#   - best parameters:
#     gamma cost
# 0.105      8
# - best performance: 0.02583333

# THIS FITS A FINAL SVM MODEL, USING THE CROSSVALIDATED TUNING PARAMETERS
# of C = (range of 8,9,10), and gamma = 0.105 or 0.11
costVal = 8
gammaVal = 0.105
svm_model <- svm(status ~., trainData, kernel = "radial",
                 cost = costVal, gamma = gammaVal)
testData$pred <- predict(svm_model, testData)
testData$correct <- with(testData, ifelse(testData$status == pred, 1, 0))
# Generate the test error
# Total number of correct classifications
totalCorrect <- sum(testData$correct)
testError_SVM <- 1 - (totalCorrect / lengthTest)
testError_SVM

## [1] 0.07692308

# Remove the "correct" column for the SVM portion
testData <- subset(testData, select = -c(correct, pred))
#===== RESULTS =====
# ALL different values here give error rate of 0.07692308, which is
# misclassifying precisely 3/39 of our test set, so 1/39 worse than for the
# linear kernel. Some of this is certainly due to there only being so much room
# to improve the error rate, since we can only pick off 1 or 2
# misclassifications in our test set.

```